



BOOK : Artificial Intelligence A Modern
Approach
Stuart J. Russell and Peter Norvig

By: Teame M.

CHAPTER ONE

Introduction to AI

2

Goals of this Course



AI is a very broad field with many subareas

- ❏ Make machines smarter (primary goal)
- ❏ Understand what intelligence is (Nobel Laureate purpose)
- ❏ Make machines more useful (entrepreneurial purpose)
- ❏ Probabilistic Learning, and Machine Learning

Artificial Intelligence



-
- Definition AI
 - Objectives /Goals of AI
 - Approaches to AI and Hypothesis of AI
 - The Foundations of AI : Bits of History and the State of the Art
 - Application Areas of AI

Artificial Intelligence



 **Artificial Intelligence made of:**

❖ **Artificial:**

- ✓ Produced by human art or effort,
- ✓ Not originating naturally.

❖ **Intelligence:**

Thinking capability

Artificial Intelligence



 What is Intelligence?

Intelligence:

- ✓ The capacity to learn and solve problems

(Websters dictionary)

In particular,

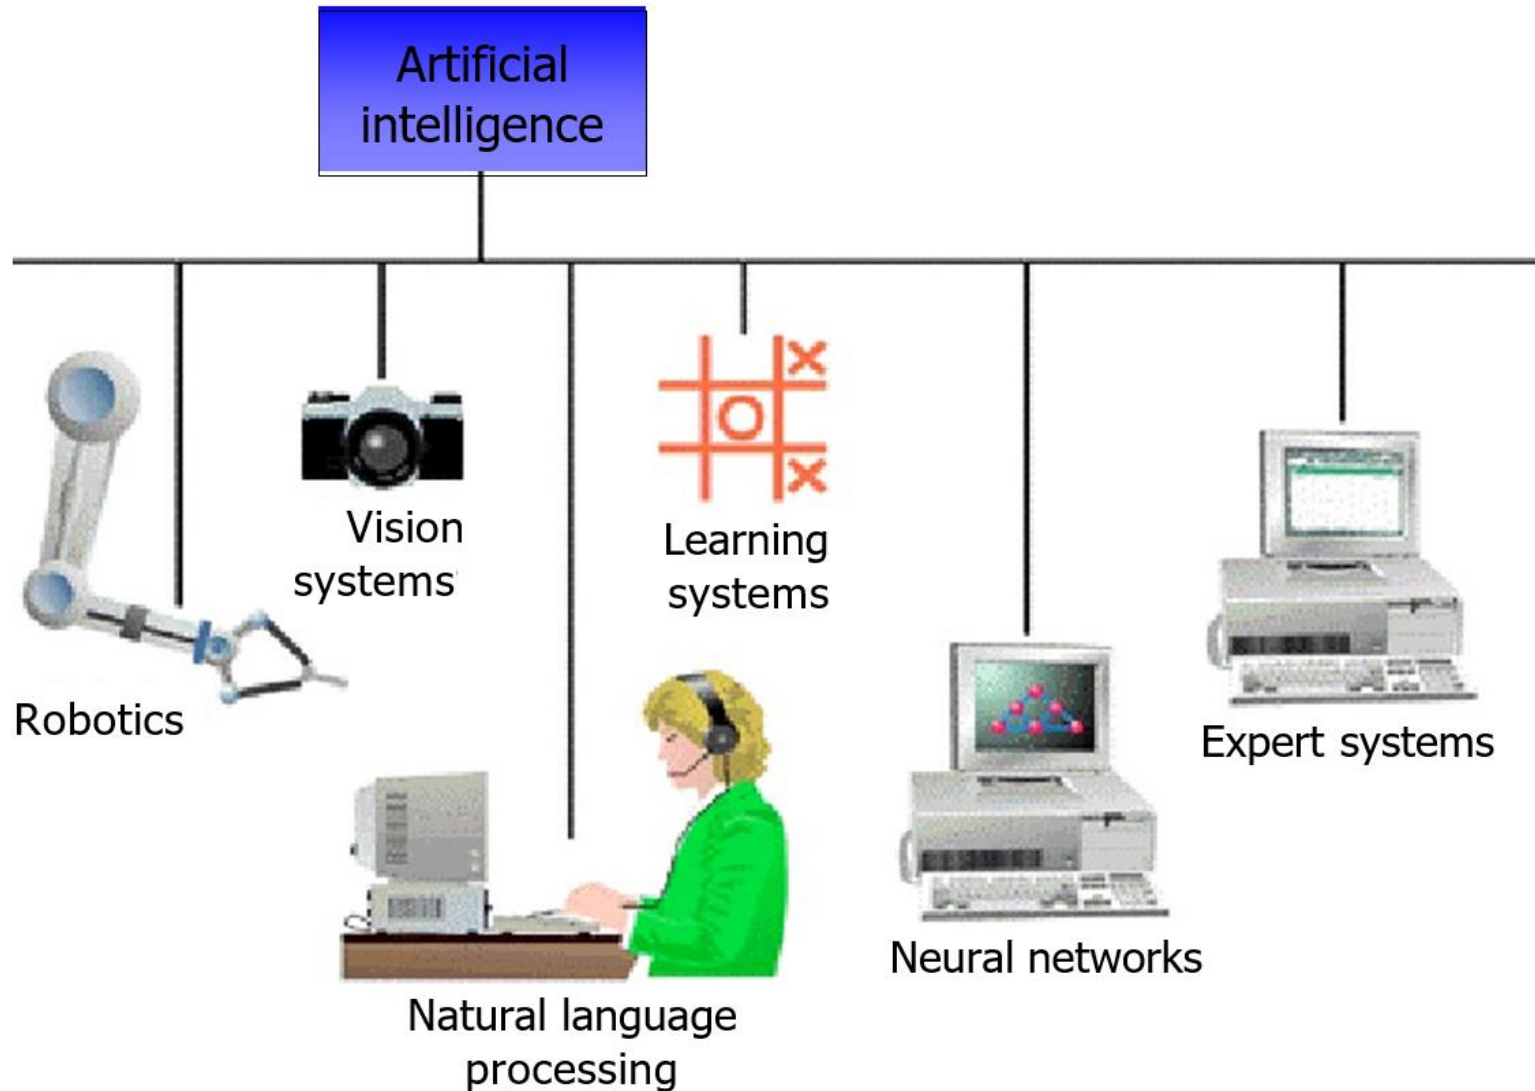
- ❑ The ability to solve novel problems
- ❑ The ability to act rationally
- ❑ The ability to act like humans
- ✓ the computational part of the ability to achieve goals in the world. Varying kinds and degrees of intelligence occur in people, many animals and some machines.

WHAT'S INVOLVED IN INTELLIGENCE?



- Ability to interact with the real world
 - to perceive, understand, and act
 - e.g., speech recognition and understanding and synthesis e.g., image understanding
- Reasoning and Planning
 - modeling the external world, given input
 - solving new problems, planning, and making decisions
- Learning and Adaptation
 - we are continuously learning and adapting
 - our internal models are always being “updated” e.g., a baby learning to categorize and recognize animals

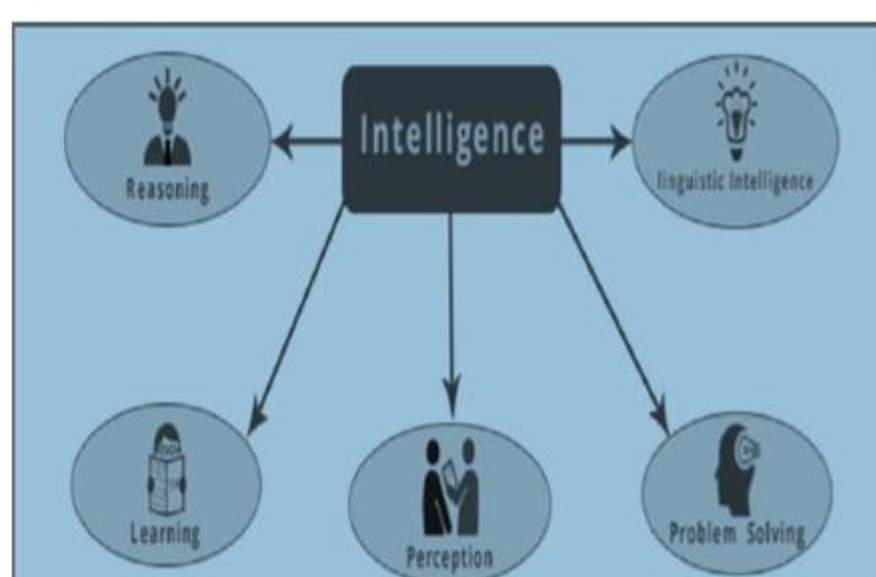
Introduction to Artificial Intelligence



Artificial Intelligence


Intelligence is composed of

- ❖ Reasoning
- ❖ Learning
- ❖ Problem Solving
- ❖ Perception
- ❖ Linguistic Intelligence



Cont...

- Let's see how raw data gets converted to intelligence through various levels of processing:



What involved in Intelligence (again)



Intelligent behavior/characteristics

- Learn from experience
- Apply knowledge acquired from experience
- Handle complex situations
- Solve problems when important information is missing
- Determine what is important
- React quickly and correctly to a new situation
- Understand visual images
- Process and manipulate symbols
- Be creative and imaginative
- Use heuristics

Objectives of AI



In general, the specific goals of AI are:


- Make machines smarter (primary goal)
- Understand what intelligence is
- Make machines more useful


Approaches of AI



- Approaches to AI do mean making computer:
- Think like a human (Thinking humanly): **The cognitive modelling approach**
 - Act like a human (Acting humanly): **The Turing Test approach**
 - Think rationally (Thinking rationally): **The "laws of thought" approach**
 - Act rationally (Acting rationally): **The rational agent approach**

Think Humanly



 If we are going to say that a given program thinks like a human, we must have some way of determining how humans think: the actual workings of human minds.


○ There are two ways to do this:

1. through introspection trying to catch our own thoughts as they go by-and
2. Through psychological experiments. Through cognitive science

This approach Requires scientific theories of internal activities of the brain.

Acting Humanly: Turing Test

- In general, the specific goals of AI are:
- Use operational qualification rather than listing intelligence qualification
- The Turing Test (1950): "Computing machinery and intelligence" for testing intelligence of machine.
- This test is proposed by Alan Turing
- He proposed computer pass that tests if a human interrogator, after posing some written questions, cannot tell whether the written responses come from a person or not Programming.



Cont..



Turing Test suggests major components required for AI

- **knowledge representation** to store what it knows.
- **Automated reasoning** :use the stored information to answer questions and to draw new conclusions;
- **computer vision**: to perceive objects
- **machine learning**: adapt to new circumstances and to detect and extrapolate patterns.
- **Natural language processing** : *to communicate successfully in NL*
- **Robotics** to manipulate objects and move about.

* Question: is it important that an intelligent system act like a human?

THINKING RATIONALLY



- ❏ Idea is from Aristotle philosophy of “**right thinking**,” that is, irrefutable reasoning processes


- ❏ His **syllogisms** provided patterns for argument structures that always yielded correct conclusions when given correct **premises**-

for example,

“Socrates is a man; all men are mortal;
therefore, Socrates is mortal.”

- ❏ These laws of thought were supposed to govern the operation of the mind; their study initiated the field called **logic**.

ACTING RATIONALLY :RATIONAL AGENT APPROACH

- 
-
- An **agent** (the Latin word to mean ager to do) is just something that acts something.
 - A **rational agent** is one that acts so as to achieve the best outcome or, when there is uncertainty, the best expected outcome.
 - Making correct inferences is sometimes part of being a rational agent,
 - All the skills needed for the Turing Test are there to allow rational actions.
 - Rational behavior
 - Doing the right thing
 - What is the “right thing”
 - That which is expected to maximize goal achievement,
 - given available information
 - We do many (“right”) things without thinking

THE FOUNDATIONS OF AI

- ❏ **Philosophy:** Logic, methods of reasoning, mind as physical system, foundations of learning, language, rationality.
- ❏ **Mathematics:** Formal representation and proof, algorithms, computation, (un)decidability, (in)tractability
- ❏ **Neuroscience:** neurons as information processing units.
- ❏ **Psychology/ :** how do people behave, perceive, process cognitive
Cognitive Science information, represent knowledge.
- ❏ **Computer:** building fast computers engineering
- ❏ **Linguistics:** knowledge representation, grammars
- ❏ **Probability/Statistics:** modeling uncertainty, learning from data
- ❏ **Economics:** utility, decision theory, rational economic agents
- ❏ **Computer Engineering:** building fast computers

History of AI








- 1943: early beginnings
- McCulloch & Pitts: Boolean circuit model of brain/ Invented Artificial neurons
- 1950: Alan Turing
- Turing's "Computing Machinery and Intelligence" was invented
- 1956: birth of AI
- Dartmouth Conference: "Artificial Intelligence" name adopted
- 1965: Robinson's complete algorithm for logical reasoning
- 1969-79: Early development of knowledge based systems took place
- 1980-- AI becomes an industry
- 1986-- Neural networks return to popularity
- 1987-- AI becomes a science
- 1995-- The emergence of intelligent agents
- Currently: robotics, drones, self-driving cars etc..

Applications of AI

 AI has been applied in different areas. Such as:

- Expert System(ES)
- Natural Language Processing(NLP)
- Speech(Voice) understanding
- Robotics and sensory
- Computer vision and scene recognition
- Intelligent Computer aided instructions
- Neural Computing
- Game Playing
- Languages Translation
- Fuzzy logic
- Genetic Algorithms
- Intelligent Agents

Home/Dorm Work

- 
-  **Can Computers Talk?**
 -  **Can Computers Recognize Speech?**
 -  **Can Computers Understand speech?**
 -  **Can Computers Learn and Adapt ?**
 -  **Can Computers “see”?**
 -  **Can computers plan and make optimal decisions?**

Can Computers Talk?



- This is known as “speech synthesis”
 - translate text to phonetic form
 - e.g., “fictitious” -> fik-tish-es
 - use pronunciation rules to map phonemes to actual sound

e.g., “tish” -> sequence of basic audio sounds

Conclusion:

- NO, for complete sentences
- YES, for individual word

Can Computers Recognize Speech?

- Speech Recognition:
 - mapping sounds from a microphone into a list of words
 - classic problem in AI, very difficult
 - “Lets talk about how to wreck a nice beach”
 - (I really said “_____”)
- Recognizing single words from a small vocabulary
- systems can do this with high accuracy (order of 99%)
- e.g., directory inquiries
 - limited vocabulary (area codes, city names)
 - computer tries to recognize you first, if unsuccessful hands you over to a human operator
 - saves millions of dollars a year for the phone companies

Recognizing human speech

- Recognizing normal speech is much more difficult
 - speech is continuous: where are the boundaries between words?
 - e.g., “John’s car has a flat tire”
 - large vocabularies
 - can be many thousands of possible words
 - we can use **context** to help figure out what someone said
 - e.g., hypothesize and test
 - try telling a waiter in a restaurant:
“I would like some cream and sugar in my coffee”
 - background noise, other speakers, accents, clogs, etc
 - on normal speech, modern systems are only about 60-70% accurate

Can Computers Understand speech?



- Understanding is different to recognition:
 - “Time flies like an arrow”
 - assume the computer can recognize all the words
 - how many different interpretations are there?
 - 1. time passes quickly like an arrow?
 - 2. command: time the flies the way an arrow times the flies
 - 3. command: only time those flies which are like an arrow
 - 4. “time-flies” are fond of arrows
 - only 1. makes any sense,
 - but how could a computer figure this out?
 - clearly humans use a lot of implicit commonsense knowledge in communication
- Conclusion: NO, much of what we say is beyond the capabilities of a computer to understand at present.

Can Computers Learn and Adapt ?

- Learning and Adaptation
 - consider a computer learning to drive on the free way
 - we could teach it lots of rules about what to do
 - or we could let it drive and steer it back on course when it heads for the embankment
 - systems like this are under development (e.g., Daimler Benz)e.g., RALPH at CMU
 - in mid 90's it drove 98% of the way from Pittsburgh to San Diego without any human assistance
 - **machine learning** allows computers to learn to do things without explicit programming.
 - many successful applications:
 - requires some “set-up”: does not mean your PC can learn to forecast the stock market or become a brain surgeon
- Conclusion: YES, computers can learn and adapt, when presented with information in the appropriate way

Can Computers “see”?

- Recognition v. Understanding (like Speech)
 - Recognition and Understanding of Objects in a scene
 - look around this room
 - you can effortlessly recognize objects
 - human brain can map 2d visual image to 3d “map”



- Why is visual recognition a hard problem?
- Conclusion:
 - mostly NO: computers can only “see” certain types of objects under limited circumstances
 - YES, for certain constrained problems (e.g., face recognition)

Can computers plan and make optimal decisions?

- Intelligence
 - involves solving problems and making decisions and plans
 - e.g., you want to take a holiday in Ethiopia
 - you need to decide on dates, flights
 - you need to get to the airport, etc
 - involves a sequence of decisions, plans, and actions
- What makes planning hard?
 - the world is not predictable:
 - your flight is canceled or there's a backup
 - there are a potentially huge number of details
 - do you consider all flights? all dates?
 - no: commonsense constrains your solutions
 - AI systems are only successful in constrained planning problems
- Conclusion: NO, real-world planning and decision-making is still beyond the capabilities of modern computers
 - exception: very well-defined, constrained problems



The End of chapter One